Invatare automata in arta vizuala

Invatare nesupervizata Retele generative

Invatare nesupervizata [1]

- Dataset
 - Doar date
 - Fare anotari
- Scop
 - o Descoperea de patternuri / structura in date.
- Exemple
 - Clusterizare
 - o Reducere dimensiunii datelor
 - o Invatarea de feature-uri
 - o Estimarea densitatii
 - Generare

"What I cannot create, I do not understand."

—Richard Feynman

Invatare nesupervizata [2] - Generare de imagini

Input

o Example din lumea reala - esantioane din distributia reala a datelor



Modelul generativ

- Genereaza exemple noi de imagini asemanatoare cu distributia pe care a fost antrenat
- O retea care genereaza la output imagini esantioane din distributia modelului

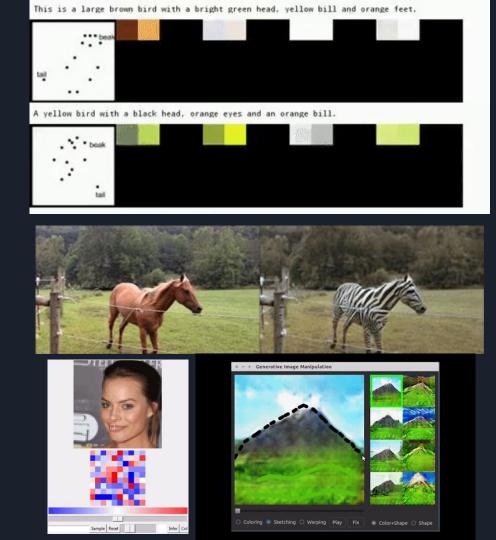
Estimarea densitatii



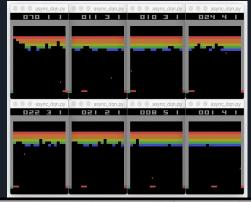
o Generare de sample-uri



- Sample-uri generate pentru arta
- Super-resolutie
- Colorizare
- Automatizarea generarii de date



Modele generative aplicate pe date time-series
 pot fi folosite pentru simulare si planificare
 (aplicatii in reinforcement learning)







Game playing

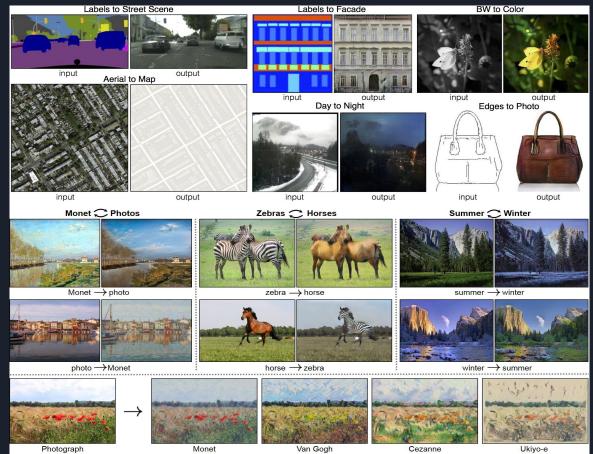


 Antrenarea modelelor generative poate determina reprezentări latente ce pot fi folosite ca feature-uri generale.

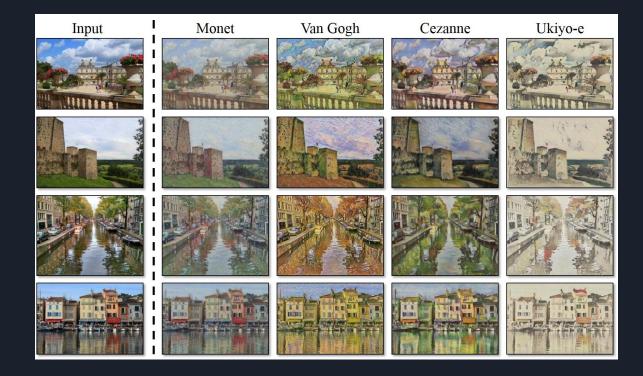




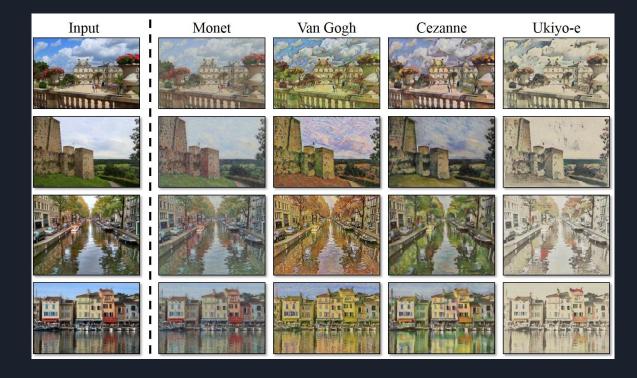
Augmentare de date



Style transfer



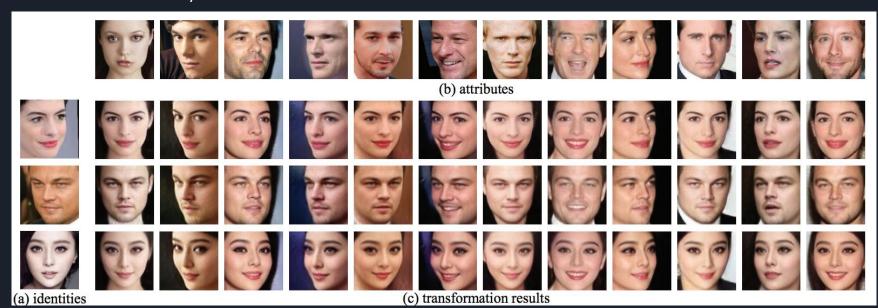
Style transfer



• Super-resolution



- Face attribute manipulation
- Face synthesis



Abordari pentru modele generative adanci

- Generative Adversarial Networks (GANs)
 - GAN
 - Wasserstein GAN
- Latent variable models
 - Variational Autoencoders (VAEs)
- Autoregressive models
 - Deep NADE
 - PixelRNN
 - PixelCNN
 - WaveNet
 - Video Pixel Network

Generative Adversarial Networks (GANs) [1]

Generator



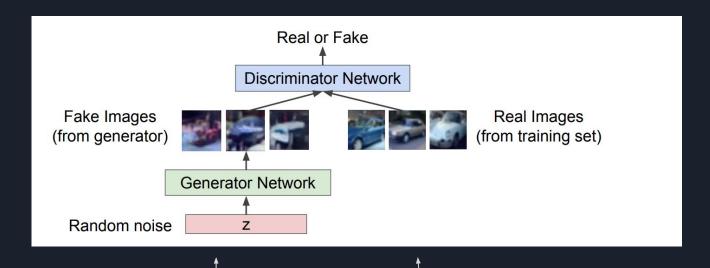
Incearca sa genereze bancnote false cat mai asemanatoare cu cele reale

Discriminator



Incearca sa distinga cat mai bine intre bancnote reale si cele false generate de catre Generator

Generative Adversarial Networks (GANs) [2]



(de obicei) generat prin samplarea unei distributii normale Datele de antrenare (training data)

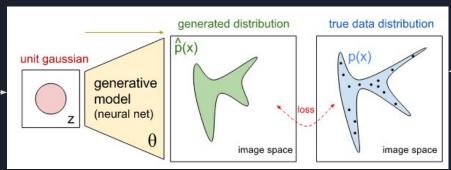
Generative Adversarial Networks (GANs) [3]

- **Generatorul** & **Discriminatorul** sunt representate de retele neuronale
- Cele doua retele sunt angrenate intr-un joc mini-max si incearca sa gaseasca un Nash Equilibrium
- Generatorul incearca sa genereze sample-uri cat mai reale astfel incat sa pacaleasca discriminatorul ca acestea sunt reale
- Discriminatorul incearca sa distinga cat mai bine intre sample-uri reale si date fake
- Antrenarea ar trebui sa se termine atunci cand
 - o Generatorul reproduce exact distributia datelor reale
 - O Discriminatorul alege random intre real/fake cu prob 0.5

Generative Adversarial Networks (GANs) [4]

Invatam o mapare (determinista) intre z si imagini generate

Z este generat prin samplarea dintr-o —— gausiana multi-dimensionala



Avem un dataset de antrenare format din exemple: x1, x2, Xn samplate din distributia reala p(x)

- Goal-ul este sa modificam parametrii retelei ai sa producem o distributie care seamana cu cea reala
- Putem folosi KL divergence ca si functie obiectiv:

$$D_{ ext{KL}}(P\|Q) = -\sum_i P(i)\,\lograc{Q(i)}{P(i)}$$

Antrenarea GAN-urilor [1]

• Jocul dintre G si D poate fi exprimat prin functia obiectiv minimax

$$\min_{G} \max_{D} \underset{\boldsymbol{x} \sim \mathbb{P}_r}{\mathbb{E}} [\log(D(\boldsymbol{x}))] + \underset{\tilde{\boldsymbol{x}} \sim \mathbb{P}_g}{\mathbb{E}} [\log(1 - D(\tilde{\boldsymbol{x}}))]$$
 Ouput discriminator pt date reale date fake

- Pr este distributia reala
- Pg este distributia modelului din care generam sample-uri:
- p(z) este o distributie simpla (normala)

$$\tilde{\boldsymbol{x}} = G(\boldsymbol{z}), \quad \boldsymbol{z} \sim p(\boldsymbol{z})$$

- Alternand
 - o Gradient ascent peste obiectivul discriminatorului

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log (1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Gradient descent peste obiectivul generatorului
 discriminator se satureaza => vanishing gradients

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

Antrenarea GAN-urilor [1]

Functia obiectiv minimax

$$\min_{G} \max_{D} \underset{\boldsymbol{x} \sim \mathbb{P}_r}{\mathbb{E}} [\log(D(\boldsymbol{x}))] + \underset{\tilde{\boldsymbol{x}} \sim \mathbb{P}_g}{\mathbb{E}} [\log(1 - D(\tilde{\boldsymbol{x}}))]$$
 Ouput discriminator pt date reale

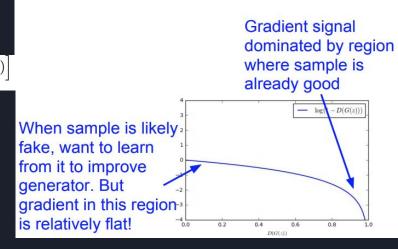
- Alternand
 - o Gradient ascent peste obiectivul discriminatorului

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log (1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

o Gradient ascent peste obiectivul generatorului

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$

In prezenta unui discriminator bun -> tot apar probleme in antrenarea generatorului



Wasserstein GAN - mai stabila antrenarea

Antrenarea GAN-urilor

for number of training iterations do

for k steps do

- Sample minibatch of m noise samples $\{z^{(1)}, \ldots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \ldots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^{m} \left[\log D_{\theta_d}(x^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(z^{(i)}))) \right]$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \ldots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by ascending its stochastic gradient (improved objective):

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(z^{(i)})))$$

end for

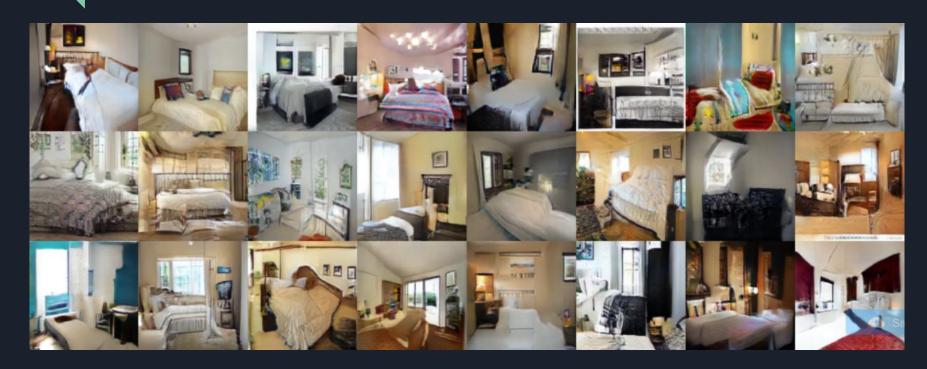
Sample-uri generate de GAN-uri





MNIST CIFAR

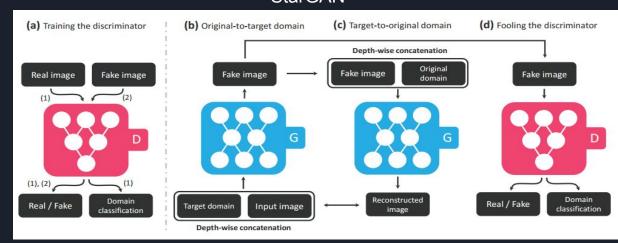
Least-Squares GAN Xudong Mao, Qing Li†, Haoran Xie, Raymond Y.K. Lau and Zhen Wang, ArXiv, Feb. 2017



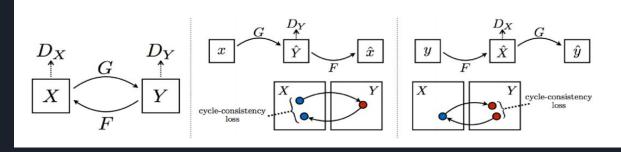
cycleGAN, StarGAN,

- Invata mapari 1-1 intre domenii folosind date ne-imperecheate (unpaired data)
- Folosesc cycle-consistency loss (reconstructie L1)

StarGAN



CycleGAN



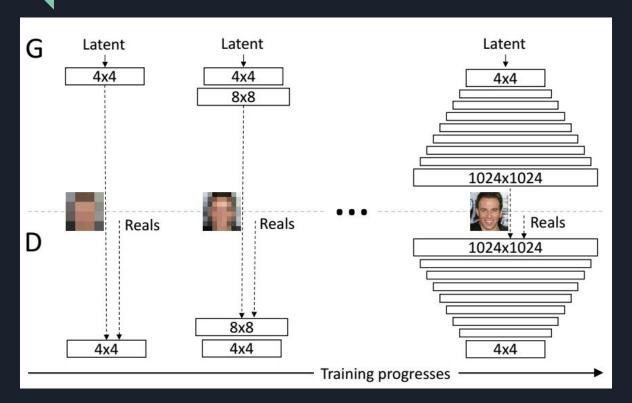
Progressive growing of GANs for improved quality, stability and variation (Kerras et al. from NVIDIA, 2017) [1]

 Imbunatateste calitatea sample-urilor crescand progresiv marimea modelului in timpul antrenarii

 Sample-urile generate cu un model antrenat pe datasetul CelebA 1024x1024

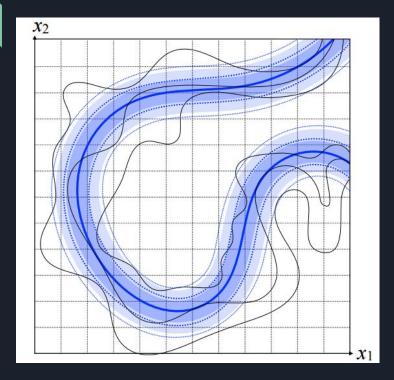


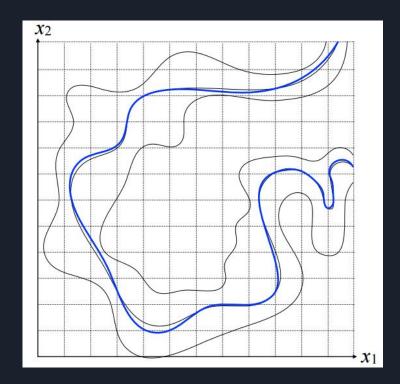
Progressive growing of GANs for improved quality, stability and variation (Kerras et al. from NVIDIA, 2017) [1]





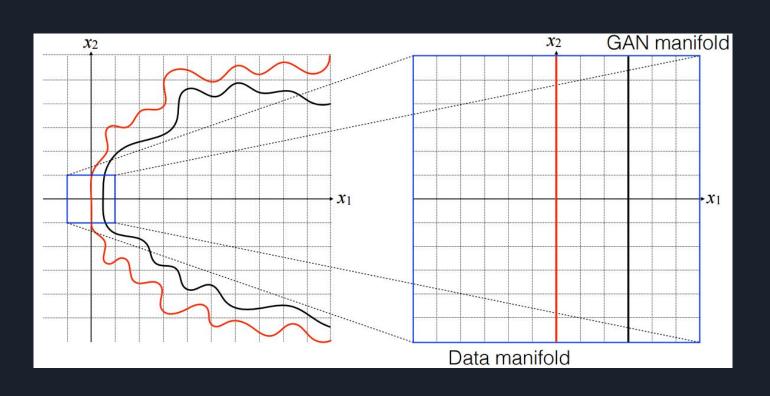
Teoria GAN-urilor





Metode traditionale - maximizarea probabilitatii de aparitie a datelor (maximum likelihood)

Antrenarea GAN-urilor: Distanta dintre manifolduri



Jensen-Shannon Divergence

 GAN-urile optimizeaza divergenta Jensen-shannon (un middle ground intre doua functii de cost)

$$JS(\mathbb{P}_r \| \mathbb{P}_g) = KL\left(\mathbb{P}_r \left\| \frac{\mathbb{P}_r + \mathbb{P}_g}{2} \right.\right) + KL\left(\mathbb{P}_g \left\| \frac{\mathbb{P}_r + \mathbb{P}_g}{2} \right.\right)$$

$$\mathrm{KL}(\mathbb{P}_r || \mathbb{P}_g) = \int \log \left(\frac{p_r(x)}{p_g(x)} \right) p_r(x) d\mu(x)$$

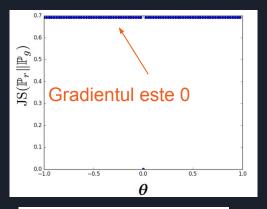
• Functia obiectiv a discriminatorului este:

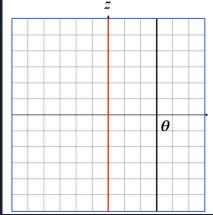
$$L(D, g_{\theta}) = \mathbb{E}_{x \sim \mathbb{P}_r}[\log D(x)] + \mathbb{E}_{x \sim \mathbb{P}_g}[\log(1 - D(x))]$$

• Un discriminator optim are forma: $D^*(x) = \frac{P_r(x)}{P_r(x) + P_q(x)}$

$$L(D^*, g_{\theta}) = 2JSD(\mathbb{P}_r || \mathbb{P}_q) - 2\log 2$$

$$JS(\mathbb{P}_r || \mathbb{P}_g) = \begin{cases} \log 2 & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0, \end{cases}$$





Distanta Earth-Movers

- Divergenta JS nu ofera semnal bun pentru antrenarea GAN-urilor
- Distanta Earth-Mover (Wasserstein-1)

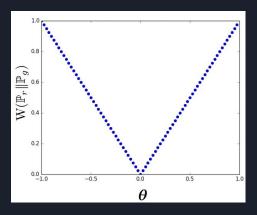
$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} \left[\|x - y\| \right]$$

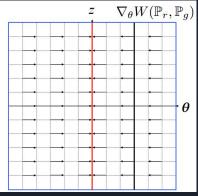
- Costum minim de a transporta masa din distributia Pr in distributia Pg
- Este continua peste tot si diferentiabila aproape peste tot (sub anumite conditii)
- Kantorovich-Rubinstein duality

$$W(\mathbb{P}_r, \mathbb{P}_g) = \sup_{\|f\|_L \le 1} \mathbb{E}_{x \sim \mathbb{P}_r}[f(x)] - \mathbb{E}_{x \sim \mathbb{P}_g}[f(x)]$$

- Supremum peste functii 1-Lipschitz f : X -> R
- o f => critic

$$W(\mathbb{P}_r || \mathbb{P}_g) = |\theta|$$



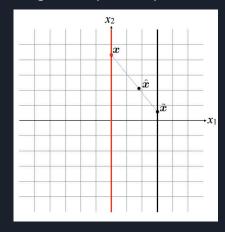


WassersteinGAN

Functia obiectiv

$$\min_{G} \max_{D \in \mathcal{D}} \ \underset{\boldsymbol{x} \sim \mathbb{P}_r}{\mathbb{E}} \left[D(\boldsymbol{x}) \right] - \underset{\tilde{\boldsymbol{x}} \sim \mathbb{P}_g}{\mathbb{E}} \left[D(\tilde{\boldsymbol{x}})) \right]$$

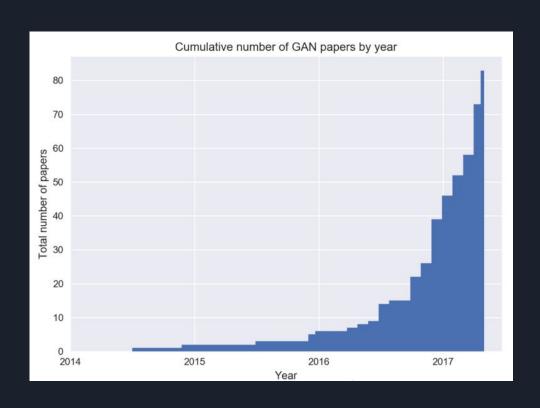
- D 1-Lipschitz functions
- Cum putem sa punem constrangerea Lipschitz peste critic D?
 - Weight clipping [-c, c]
 - Gradient penalty



$$\mathbb{E}_{\hat{\boldsymbol{x}} \sim \mathbb{P}_{\hat{\boldsymbol{x}}}} \left[(\|\nabla_{\hat{\boldsymbol{x}}} D(\hat{\boldsymbol{x}})\|_2 - 1)^2 \right]$$

$$egin{aligned} \epsilon \sim U[0,1], oldsymbol{x} \sim \mathbb{P}_r, & ilde{oldsymbol{x}} \sim \mathbb{P}_g \ & \hat{oldsymbol{x}} = \epsilon oldsymbol{x} + (1-\epsilon) & ilde{oldsymbol{x}} \end{aligned}$$

Explozia GAN-urilor



References

- https://drive.google.com/file/d/0ByUKRdiCDK7-bTgxTGoxYjQ4NW8/view?usp=drive web - Generative Models 1
- https://drive.google.com/file/d/0B wzP JIVFcKQ21udGpTSkh0aVk/view?usp=drive we
 b Generative Models 2
- https://arxiv.org/pdf/1701.04862.pdf GAN theory
- https://arxiv.org/abs/1701.07875 WassersteinGAN
- http://arxiv.org/pdf/1711.09020v1.pdf StarGAN
- https://arxiv.org/abs/1703.10593 CycleGAN
- https://arxiv.org/abs/1710.10196v3 Progressive Growing of GANs for Improved Quality, Stability, and Variation
- https://arxiv.org/abs/1511.06434 DCGAN
- https://arxiv.org/abs/1701.00160 NIPS GAN tutorial

